

Vector Calculus

Last modified: 2026-10-09 11:00:09 (PDT)

(adapted from Fieller (2016), [Section 7.2](#))

This section covers [derivatives](#) of functions of vectors and matrices. Its linear algebra prerequisites, such as vectors, matrices, transposes, dot products, and quadratic forms, are covered in [Linear Algebra](#).

Let $\tilde{x}$ and $\tilde{\beta}$ be column vectors of length p (see [column vector](#) and [dot product](#)).

Definition 0.1 (Partial derivative). Let $f : \mathbb{R}^p \rightarrow \mathbb{R}$ and $\tilde{x} = (x_1, \dots, x_p)^\top \in \mathbb{R}^p$. The **partial derivative** of f with respect to x_j at $\tilde{x}$ is

$$\frac{\partial}{\partial x_j} f(\tilde{x}) \stackrel{\text{def}}{=} \lim_{h \rightarrow 0} \frac{f(x_1, \dots, x_j + h, \dots, x_p) - f(x_1, \dots, x_p)}{h},$$

when this [limit](#) exists. It is the [derivative](#) at x_j of the one-variable function $t \mapsto f(x_1, \dots, t, \dots, x_p)$, which varies the j th coordinate and holds the other coordinates fixed.

Example 0.1 (Partial derivatives of $x_1^2 x_2 + 3x_2$). Let $f(x_1, x_2) = x_1^2 x_2 + 3x_2$. At $\tilde{x} = (1, 2)^\top$,

$$\begin{aligned} f(1, 2) &= 1 \cdot 2 + 6 \\ &= 8, \end{aligned}$$

and for $h \neq 0$

$$\begin{aligned}
\frac{f(1+h, 2) - f(1, 2)}{h} &= \frac{(1+h)^2 \cdot 2 + 6 - 8}{h} && \text{(substitute into } f) \\
&= \frac{(1+h)^2 \cdot 2 - 2}{h} && (6 - 8 = -2) \\
&= \frac{(1 + 2h + h^2) \cdot 2 - 2}{h} && \text{(expand } (1+h)^2) \\
&= \frac{2 + 4h + 2h^2 - 2}{h} && \text{(distributive law)} \\
&= \frac{4h + 2h^2 + 2 - 2}{h} && \text{(reorder the terms)} \\
&= \frac{4h + 2h^2}{h} && (2 - 2 = 0) \\
&= 4 + 2h, && \text{(divide by } h \neq 0)
\end{aligned}$$

which tends to 4 as $h \rightarrow 0$, so $\frac{\partial}{\partial x_1} f(1, 2) = 4$. Holding x_2 fixed and differentiating in x_1 gives the same answer: $\frac{\partial}{\partial x_1} f(\tilde{x}) = 2x_1 x_2$, which is $2 \cdot 1 \cdot 2 = 4$ at $(1, 2)^\top$. Holding x_1 fixed instead gives $\frac{\partial}{\partial x_2} f(\tilde{x}) = x_1^2 + 3$, which is $1 + 3 = 4$ at $(1, 2)^\top$.

Exercise 0.1 (First-order partial derivatives). Adapted from Miller (2016), Question 1.1.42.

Let $f(x, y) = x^2 y + e^x + \sin(xy)$. Find $\frac{\partial f}{\partial x}$ and $\frac{\partial f}{\partial y}$.

Solution

Solution 0.1. To compute $\frac{\partial f}{\partial x}$, treat y as a constant:

$$\frac{\partial f}{\partial x} = 2xy + e^x + y \cos(xy)$$

To compute $\frac{\partial f}{\partial y}$, treat x as a constant:

$$\begin{aligned}
\frac{\partial f}{\partial y} &= x^2 + 0 + x \cos(xy) \\
&= x^2 + x \cos(xy)
\end{aligned}$$

Exercise 0.2 (Partial derivatives of the normal density with respect to parameters). Adapted from Miller (2016), Question 1.1.43.

Let

$$f(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma^2} \exp\left\{-\frac{(x-\mu)^2}{2\sigma^2}\right\}$$

Find $\frac{\partial f}{\partial \mu}$ and $\frac{\partial f}{\partial \sigma}$.

Solution

Solution 0.2. 1. Derivative with respect to μ : Treat σ as constant. Differentiating the exponential factor via the chain rule:

$$\begin{aligned} \frac{\partial f}{\partial \mu} &= \frac{1}{\sqrt{2\pi}\sigma^2} \exp\left\{-\frac{(x-\mu)^2}{2\sigma^2}\right\} \cdot \left[-\frac{2(x-\mu)(-1)}{2\sigma^2}\right] \\ &= \frac{x-\mu}{\sigma^2} f(x; \mu, \sigma) \end{aligned}$$

2. Derivative with respect to σ : Write $f(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi}}\sigma^{-1} \exp\left\{-\frac{(x-\mu)^2}{2\sigma^2}\right\}$ and apply the product rule:

$$\begin{aligned} \frac{\partial f}{\partial \sigma} &= -\frac{1}{\sqrt{2\pi}}\sigma^{-2} \exp\left\{-\frac{(x-\mu)^2}{2\sigma^2}\right\} + \frac{1}{\sqrt{2\pi}}\sigma^{-1} \exp\left\{-\frac{(x-\mu)^2}{2\sigma^2}\right\} \cdot \left[\frac{(x-\mu)^2}{\sigma^3}\right] \\ &= \left[-\frac{1}{\sigma} + \frac{(x-\mu)^2}{\sigma^3}\right] f(x; \mu, \sigma) \end{aligned}$$

Remark: Setting $\frac{\partial f}{\partial \mu} = 0$ and $\frac{\partial f}{\partial \sigma} = 0$ for a sample of independent observations leads directly to the maximum likelihood estimators $\hat{\mu} = \bar{x}$ and $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$.

Exercise 0.3 (Partial derivatives with exponential of quadratic form). Adapted from Miller (2016), Question 1.1.44.

Find $\frac{\partial f}{\partial x}$ and $\frac{\partial f}{\partial y}$ for

$$f(x, y) = x \exp\{x^2 + y^2\}$$

Solution

Solution 0.3. With respect to x , apply the product rule:

$$\begin{aligned} \frac{\partial f}{\partial x} &= (1) \exp\{x^2 + y^2\} + x \cdot [2x \exp\{x^2 + y^2\}] \\ &= (1 + 2x^2) \exp\{x^2 + y^2\} \end{aligned}$$

With respect to y , treat x as constant:

$$\begin{aligned}\frac{\partial f}{\partial y} &= x \cdot [2y \exp\{x^2 + y^2\}] \\ &= 2xy \exp\{x^2 + y^2\}\end{aligned}$$

Exercise 0.4 (Using symmetry to compute partial derivatives). Adapted from Miller (2016), Question 1.1.45.
Find $\frac{\partial f}{\partial x}$ and $\frac{\partial f}{\partial y}$ for

$$f(x, y) = e^{xy} - \log(x^2 + y^2)$$

on $\mathbb{R}^2 \setminus \{(0, 0)\}$.

Solution

Solution 0.4. Holding y constant, differentiate with respect to x :

$$\frac{\partial f}{\partial x} = ye^{xy} - \frac{2x}{x^2 + y^2}$$

Notice that the function is symmetric in x and y : $f(x, y) = f(y, x)$. Interchanging x and y gives $\frac{\partial f}{\partial y}$ immediately:

$$\frac{\partial f}{\partial y} = xe^{xy} - \frac{2y}{x^2 + y^2}$$

Remark: Recognizing symmetry in multi-variable functions eliminates redundant calculations.

Definition 0.2 (Vector derivative (gradient)). If $f(\tilde{\beta})$ is a scalar-valued function of a $p \times 1$ vector $\tilde{\beta}$, such as $f(\tilde{\beta}) = \tilde{x} \cdot \tilde{\beta}$, then its **vector derivative** is the column vector of its partial derivatives (Definition 0.1):

$$\frac{\partial}{\partial \tilde{\beta}} f(\tilde{\beta}) = \begin{bmatrix} \frac{\partial}{\partial \beta_1} f(\tilde{\beta}) \\ \frac{\partial}{\partial \beta_2} f(\tilde{\beta}) \\ \vdots \\ \frac{\partial}{\partial \beta_p} f(\tilde{\beta}) \end{bmatrix}$$

The vector derivative is also called the **gradient** of f with respect to $\tilde{\beta}$, written $\nabla_{\tilde{\beta}} f(\tilde{\beta})$.

Example 0.2 (A vector derivative). Let $\tilde{\beta} = (\beta_1, \beta_2)^\top$ and $f(\tilde{\beta}) = 3\beta_1 + 5\beta_2^2$. Then

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}}_{2 \times 1} f(\tilde{\beta}) = \begin{bmatrix} \frac{\partial}{\partial \beta_1} (3\beta_1 + 5\beta_2^2) \\ \frac{\partial}{\partial \beta_2} (3\beta_1 + 5\beta_2^2) \end{bmatrix} = \begin{bmatrix} 3 \\ 10\beta_2 \end{bmatrix},$$

and at $\tilde{\beta} = (1, 2)^\top$ the vector derivative is $(3, 20)^\top$.

💡 Video lecture

Hutchinson’s [Gradients Refresher](#) (17 min) covers gradients and partial derivatives, the ideas behind this section (Hutchinson, n.d.). The login for the video site is posted [on Canvas](#).

💡 Videos

Jon Krohn’s “Calculus for Machine Learning” YouTube playlist has videos on partial derivatives:

- [What Partial Derivatives Are \(Hands-on Introduction\)](#)
- [Partial Derivative Exercises](#)
- [Advanced Partial Derivatives](#)
- [Advanced Partial-Derivative Exercises](#)
- [Partial Derivative Notation](#)

Checking a gradient by its shape

The gradient always has the same shape as the input it is taken with respect to, because it holds exactly one partial derivative per component of that input:

Table 1: The shape of a gradient follows the shape of what it differentiates with respect to.

f is a function of	its gradient is
a single number	a single number — the ordinary derivative
a vector in $\mathbb{R}^p$	a vector in $\mathbb{R}^p$
a matrix in $\mathbb{R}^{m \times n}$	a matrix in $\mathbb{R}^{m \times n}$

The table is the cheapest check there is on a gradient just worked out by hand: if the shapes

disagree, something went wrong. It is the shape rule of matrix multiplication again, applied to calculus.

Definition 0.3 (Direction). A **direction** in $\mathbb{R}^p$ is given by a nonzero vector $\tilde{d} \in \mathbb{R}^p$: moving from a point $\tilde{x}$ in the direction $\tilde{d}$ means moving to the points $\tilde{x} + t\tilde{d}$ with $t > 0$. For any $c > 0$, $c\tilde{d}$ reaches the same points (with t/c in place of t), so $\tilde{d}$ and $c\tilde{d}$ give the same direction.

Example 0.3 (Moving from $(2, -1)^\top$ in the direction $(1, 1)^\top$). From $\tilde{x} = (2, -1)^\top$, moving in the direction $\tilde{d} = (1, 1)^\top$ reaches the points $\tilde{x} + t\tilde{d} = (2 + t, -1 + t)^\top$ for $t > 0$, such as $(3, 0)^\top$ at $t = 1$. The vector $(2, 2)^\top$ gives the same direction: it reaches $(3, 0)^\top$ at $t = \frac{1}{2}$. The vector $(-1, -1)^\top$ gives the opposite direction, reaching points such as $(1, -2)^\top$.

Definition 0.4 (Level set (level curve, contour)). Let $f : \mathbb{R}^p \rightarrow \mathbb{R}$ and let c be a real number. The **level set** of f at c is

$$\{\tilde{x} \in \mathbb{R}^p : f(\tilde{x}) = c\},$$

the set of points where f takes the value c . For $p = 2$, a level set is often a curve in the plane, called a **level curve** (or **contour**) of f .

Example 0.4 (Level sets of $x_1^2 + x_2^2$). Let $f(\tilde{x}) = x_1^2 + x_2^2$ for $\tilde{x} = (x_1, x_2)^\top$.

- The level set at 1 is the circle of radius 1 around $\tilde{0}$. For example, $(0.6, 0.8)^\top$ is on it, since $0.36 + 0.64 = 1$, and $(1, 1)^\top$ is not, since $1 + 1 = 2$.
- The level set at 0 is the single point $\tilde{0}$.
- The level set at -1 is empty, since $x_1^2 + x_2^2 \geq 0$.

The course website's version of this page has an interactive plot [here](#).

Exercise 0.5 (Take a gradient by hand). Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ be

$$f(\tilde{w}) = 3w_1w_2 + w_2^3 - 5w_1$$

1. Compute $\partial f / \partial w_1$ (Definition 0.1), treating w_2 as a constant.
2. Compute $\partial f / \partial w_2$, treating w_1 as a constant.
3. Assemble $\nabla_{\tilde{w}} f$ and evaluate it at $\tilde{w} = [2 \ -1]^\top$.
4. Check the shape of your answer against Table 1.

Solution

Solution 0.5. 1. Differentiate with respect to w_1 , holding w_2 fixed. The first term is the constant $3w_2$ times w_1 , so it contributes $3w_2$. The second term has no w_1 in it, so it contributes 0. The third term is -5 times w_1 , so it contributes -5 :

$$\begin{aligned}\frac{\partial f}{\partial w_1} &= 3w_2 + 0 - 5 \\ &= 3w_2 - 5\end{aligned}$$

2. Now with respect to w_2 , holding w_1 fixed. The first term is the constant $3w_1$ times w_2 , contributing $3w_1$. The second term contributes $3w_2^2$. The third has no w_2 in it, contributing 0:

$$\begin{aligned}\frac{\partial f}{\partial w_2} &= 3w_1 + 3w_2^2 + 0 \\ &= 3w_1 + 3w_2^2\end{aligned}$$

3. Stack the two, in the order the coordinates are numbered:

$$\nabla_{\tilde{w}} f(\tilde{w}) = \begin{bmatrix} 3w_2 - 5 \\ 3w_1 + 3w_2^2 \end{bmatrix}$$

The gradient is a vector-valued *function* of $\tilde{w}$, not a single vector. At $\tilde{w} = \begin{bmatrix} 2 & -1 \end{bmatrix}^\top$,

$$\begin{aligned}\nabla_{\tilde{w}} f &= \begin{bmatrix} 3(-1) - 5 \\ 3(2) + 3(-1)^2 \end{bmatrix} \\ &= \begin{bmatrix} -3 - 5 \\ 6 + 3 \end{bmatrix} \\ &= \begin{bmatrix} -8 \\ 9 \end{bmatrix}\end{aligned}$$

4. The input was a vector in $\mathbb{R}^2$ and so is the answer, as Table 1 requires. Note the signs: from this point, increasing w_1 *decreases* f while increasing w_2 increases it, so the uphill direction is neither axis.

Definition 0.5 (Row-vector derivative (total derivative)). If $f(\tilde{\beta})$ is a scalar-valued function of a $p \times 1$ vector $\tilde{\beta}$, such as $f(\tilde{\beta}) = \tilde{x} \cdot \tilde{\beta}$, then its **row-vector derivative** is the row vector of its partial derivatives (Definition 0.1):

$$\frac{\partial f(\tilde{\beta})}{\partial \tilde{\beta}^\top} = \begin{bmatrix} \frac{\partial}{\partial \beta_1} f(\tilde{\beta}) & \frac{\partial}{\partial \beta_2} f(\tilde{\beta}) & \cdots & \frac{\partial}{\partial \beta_p} f(\tilde{\beta}) \end{bmatrix} \quad (1)$$

Some sources write the same row vector as $\frac{\partial}{\partial \tilde{\beta}^\top} f(\tilde{\beta})$, with the operator on the left, read as each operator $\frac{\partial}{\partial \beta_i}$ applied to f rather than as a matrix product (see Remark 0.1). The row-vector derivative is also called the **total derivative** of f , the name used when the derivative of a scalar-valued function is written as a row vector.

Remark 0.1 (Which side the operator goes on). Read $\frac{\partial}{\partial \tilde{\beta}}$ as a $p \times 1$ column vector of operators with entries $\frac{\partial}{\partial \beta_1}, \dots, \frac{\partial}{\partial \beta_p}$, and $\frac{\partial}{\partial \tilde{\beta}^\top}$ as the $1 \times p$ row vector of operators with the same entries. Writing either one next to f can be read in two ways:

- **As a matrix product**, with f a 1×1 matrix, under the shape rule of [matrix multiplication](#). The column-vector operator works on the left: $\frac{\partial}{\partial \tilde{\beta}} f$ is $(p \times 1)(1 \times 1)$, a $p \times 1$ column vector (Definition 0.2). The row-vector operator does not: $\frac{\partial}{\partial \tilde{\beta}^\top} f$ is $(1 \times p)(1 \times 1)$, whose inner dimensions p and 1 do not match when $p > 1$. For the shape rule to give the $1 \times p$ row vector, the operator would have to go on the right of f , as in $(1 \times 1)(1 \times p)$, which changes its meaning (see below).
- **As a scalar multiple**, with f a scalar, the way $c\tilde{x}$ is read in [scalar multiplication](#), which has no shape rule: each entry of the operator vector is applied to f . Then both $\frac{\partial}{\partial \tilde{\beta}} f$ and $\frac{\partial}{\partial \tilde{\beta}^\top} f$ work with the operator on the left.

The fraction $\frac{\partial f}{\partial \tilde{\beta}^\top}$ of Equation 1 means the same thing under both readings: the transpose in the denominator says the result is a row vector, without making the derivative a product.

Treating the operators as vectors is sound for checking shapes, but “multiplying” an operator by a function means applying the operator to it, so two rules of ordinary matrix algebra do not carry over: scalars commute with matrices, and products can be regrouped.

- **Scalars do not commute with operators.** $\frac{\partial}{\partial \beta_i} f$ is a function, but $f \frac{\partial}{\partial \beta_i}$ is still an unapplied operator, waiting for something to act on. So, unlike $c\tilde{x} = \tilde{x}c$, the operator cannot move to the other side of f without changing the meaning. Putting the row-vector operator on the right of f gives the row-vector derivative only under the convention that it acts on the factor to its left.
- **Products cannot be regrouped.** An operator acts on the whole product to its right, so by the [product rule](#), $\frac{\partial}{\partial \beta_i}(fg) \neq \left(\frac{\partial}{\partial \beta_i} f\right)g$ in general. For example, with $f(\tilde{\beta}) = \beta_1$, $\frac{\partial}{\partial \beta_1}(ff) = 2\beta_1$ but $\left(\frac{\partial}{\partial \beta_1} f\right)f = \beta_1$.

Theorem 0.1 (Row and column derivatives are transposes).

$$\frac{\partial f(\tilde{\beta})}{\partial \tilde{\beta}^\top} = \left(\frac{\partial f(\tilde{\beta})}{\partial \tilde{\beta}} \right)^\top \quad (2)$$

$$\frac{\partial f(\tilde{\beta})}{\partial \tilde{\beta}} = \left(\frac{\partial f(\tilde{\beta})}{\partial \tilde{\beta}^\top} \right)^\top \quad (3)$$

i Proof

Proof. By Definition 0.2 and Definition 0.5, entry j of both $\frac{\partial f(\tilde{\beta})}{\partial \tilde{\beta}}$ and $\frac{\partial f(\tilde{\beta})}{\partial \tilde{\beta}^\top}$ is $\frac{\partial}{\partial \beta_j} f(\tilde{\beta})$; the first is a $p \times 1$ column vector and the second a $1 \times p$ row vector with the same entries in the same order, so each is the transpose of the other. $\square$

Example 0.5 (Row and column derivatives of a linear function). For $f(\tilde{\beta}) = 3\beta_1 + 5\beta_2$:

$$\frac{\partial f(\tilde{\beta})}{\partial \tilde{\beta}} = \begin{bmatrix} 3 \\ 5 \end{bmatrix}, \quad \frac{\partial f(\tilde{\beta})}{\partial \tilde{\beta}^\top} = [3 \quad 5],$$

and each is the transpose of the other.

Definition 0.6 (Derivative of a vector-valued function). If

$$\begin{aligned} \tilde{y} &= \tilde{y}(\tilde{\beta}) \\ &= (y_1, \dots, y_q)^\top \end{aligned}$$

is a $q \times 1$ vector-valued function of the $p \times 1$ vector $\tilde{\beta}$, its **derivative with respect to $\tilde{\beta}$** is the $p \times q$ matrix whose (i, j) entry is

$$\left[\frac{\partial}{\partial \tilde{\beta}} \tilde{y}^\top \right]_{ij} \stackrel{\text{def}}{=} \frac{\partial}{\partial \beta_i} y_j, \quad i = 1, \dots, p, \quad j = 1, \dots, q.$$

Writing derivatives this way is the **denominator layout**, which these notes use throughout: rows index the entries of $\tilde{\beta}$ (the denominator) and columns index the entries of $\tilde{y}$ (the numerator), so column j is the vector derivative $\frac{\partial}{\partial \tilde{\beta}} y_j$ (Definition 0.2). Both $\frac{\partial}{\partial \tilde{\beta}} \tilde{y}$ and $\frac{\partial}{\partial \tilde{\beta}} \tilde{y}^\top$ denote this $p \times q$ matrix.

Example 0.6 (Differentiating a 3×1 function of a 2×1 vector). Let $\tilde{\beta} = (\beta_1, \beta_2)^\top$ ($p = 2$) and $\tilde{y}(\tilde{\beta}) = (\beta_1^2, \beta_1\beta_2, 3\beta_2)^\top$ ($q = 3$). Then

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}} \tilde{y}^\top}_{2 \times 3} = \begin{bmatrix} \frac{\partial}{\partial \beta_1} \beta_1^2 & \frac{\partial}{\partial \beta_1} \beta_1 \beta_2 & \frac{\partial}{\partial \beta_1} 3\beta_2 \\ \frac{\partial}{\partial \beta_2} \beta_1^2 & \frac{\partial}{\partial \beta_2} \beta_1 \beta_2 & \frac{\partial}{\partial \beta_2} 3\beta_2 \end{bmatrix} \\ = \begin{bmatrix} 2\beta_1 & \beta_2 & 0 \\ 0 & \beta_1 & 3 \end{bmatrix}$$

Definition 0.7 (Jacobian matrix). Let

$$\tilde{y} = \tilde{y}(\tilde{\beta}) \\ = (y_1, \dots, y_q)^\top$$

be a $q \times 1$ vector-valued function of the $p \times 1$ vector $\tilde{\beta}$. The **Jacobian matrix** of $\tilde{y}$ is the $q \times p$ matrix whose (j, i) entry is $\frac{\partial}{\partial \beta_i} y_j$, for $j = 1, \dots, q$ and $i = 1, \dots, p$. It is the transpose of the derivative $\frac{\partial}{\partial \tilde{\beta}} \tilde{y}^\top$ of Definition 0.6. Writing derivatives this way, with rows indexing the entries of $\tilde{y}$ (the numerator) and columns indexing the entries of $\tilde{\beta}$ (the denominator), is the **numerator layout**.

Example 0.7 (The Jacobian matrix of a 3×1 function of a 2×1 vector). For $\tilde{y}(\tilde{\beta}) = (\beta_1^2, \beta_1\beta_2, 3\beta_2)^\top$ of Example 0.6, the Jacobian matrix is the 3×2 matrix

$$\begin{bmatrix} \frac{\partial}{\partial \beta_1} \beta_1^2 & \frac{\partial}{\partial \beta_2} \beta_1^2 \\ \frac{\partial}{\partial \beta_1} \beta_1 \beta_2 & \frac{\partial}{\partial \beta_2} \beta_1 \beta_2 \\ \frac{\partial}{\partial \beta_1} 3\beta_2 & \frac{\partial}{\partial \beta_2} 3\beta_2 \end{bmatrix} = \begin{bmatrix} 2\beta_1 & 0 \\ \beta_2 & \beta_1 \\ 0 & 3 \end{bmatrix},$$

the transpose of the 2×3 matrix found there.

Remark 0.2 (Numerator layout). Some sources use the numerator layout (Definition 0.7) throughout, so that their derivative of a vector-valued function is the Jacobian matrix, the transpose of the derivative used in these notes. In numerator layout, the derivative of a scalar-valued function is a row vector: the row-vector derivative (Definition 0.5). Check a source's layout before combining its formulas with these.

Definition 0.8 (Constant). A $q \times 1$ vector $\tilde{x}$ is **constant with respect to** the $p \times 1$ vector $\tilde{\beta}$ if its derivative (Definition 0.6) is zero:

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}} \tilde{x}^\top}_{p \times q} = \underset{p \times q}{\mathbf{0}}$$

Example 0.8 (A constant vector). Let $\tilde{\beta} = (\beta_1, \beta_2)^\top$ and $\tilde{x} = (3, 5)^\top$, so $x_1 = 3$ and $x_2 = 5$ do not depend on $\tilde{\beta}$. Expanding $\frac{\partial}{\partial \tilde{\beta}} \tilde{x}^\top$ into its matrix of scalar partial derivatives (Definition 0.6) and evaluating each entry:

$$\begin{aligned} \underbrace{\frac{\partial}{\partial \tilde{\beta}} \tilde{x}^\top}_{2 \times 2} &= \frac{\partial}{\partial \tilde{\beta}} \begin{bmatrix} x_1 & x_2 \end{bmatrix} \\ &= \begin{bmatrix} \frac{\partial}{\partial \beta_1} x_1 & \frac{\partial}{\partial \beta_1} x_2 \\ \frac{\partial}{\partial \beta_2} x_1 & \frac{\partial}{\partial \beta_2} x_2 \end{bmatrix} \\ &= \begin{bmatrix} \frac{\partial}{\partial \beta_1} 3 & \frac{\partial}{\partial \beta_1} 5 \\ \frac{\partial}{\partial \beta_2} 3 & \frac{\partial}{\partial \beta_2} 5 \end{bmatrix} \\ &= \begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix} \\ &= \underset{2 \times 2}{\mathbf{0}} \end{aligned}$$

Every entry is the derivative of a constant, so $\frac{\partial}{\partial \tilde{\beta}} \tilde{x}^\top = \underset{2 \times 2}{\mathbf{0}}$ and $\tilde{x}$ is constant with respect to $\tilde{\beta}$ (Definition 0.8).

Example 0.9 (A vector that is not constant). With the same $\tilde{\beta}$, let $\tilde{x} = (\beta_1, 3)^\top$. Then

$$\begin{aligned} \underbrace{\frac{\partial}{\partial \tilde{\beta}} \tilde{x}^\top}_{2 \times 2} &= \begin{bmatrix} \frac{\partial}{\partial \beta_1} \beta_1 & \frac{\partial}{\partial \beta_1} 3 \\ \frac{\partial}{\partial \beta_2} \beta_1 & \frac{\partial}{\partial \beta_2} 3 \end{bmatrix} \\ &= \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}, \end{aligned}$$

which is not $\underset{2 \times 2}{\mathbf{0}}$, so this $\tilde{x}$ is not constant with respect to $\tilde{\beta}$: its first entry changes when β_1 does.

Exercise 0.6 (Partial derivative with respect to an unreferenced variable). Adapted from Miller (2016), Question 1.1.46.
Find $\frac{\partial f}{\partial x}$ and $\frac{\partial f}{\partial y}$ for

$$f(x, y, t) = 5t^4 - 4t^5 \cos(t \sin t)$$

Solution

Solution 0.6. Although the function is declared as a function of three variables (x, y, t) , its expression depends solely on t .

Because neither x nor y appears in the formula, f is constant with respect to both x and y :

$$\frac{\partial f}{\partial x} = 0, \quad \frac{\partial f}{\partial y} = 0$$

Remark: Always verify which variables actually appear in a formula before performing lengthy algebraic differentiation.

Theorem 0.2 (Derivative of a dot product). *If $\tilde{x}$ is constant with respect to $\tilde{\beta}$, then:*

$$\begin{aligned} \underbrace{\frac{\partial}{\partial \tilde{\beta}}(\tilde{x} \cdot \tilde{\beta})}_{p \times 1} &= \underbrace{\frac{\partial}{\partial \tilde{\beta}}(\tilde{\beta} \cdot \tilde{x})}_{p \times 1} \\ &= \underbrace{\tilde{x}}_{p \times 1} \end{aligned}$$

Proof

Proof.

$$\begin{aligned} \frac{\partial}{\partial \tilde{\beta}}(\tilde{x} \cdot \tilde{\beta}) &= \begin{bmatrix} \frac{\partial}{\partial \beta_1}(x_1\beta_1 + x_2\beta_2 + \dots + x_p\beta_p) \\ \frac{\partial}{\partial \beta_2}(x_1\beta_1 + x_2\beta_2 + \dots + x_p\beta_p) \\ \vdots \\ \frac{\partial}{\partial \beta_p}(x_1\beta_1 + x_2\beta_2 + \dots + x_p\beta_p) \end{bmatrix} \\ &= \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_p \end{bmatrix} \\ &= \tilde{x} \end{aligned}$$

□

Example 0.10 (Derivative of a dot product). Let $\tilde{x} = (3, 5)^\top$ (constant with respect to $\tilde{\beta}$; see Example 0.8) and $\tilde{\beta} = (\beta_1, \beta_2)^\top$. Then $\tilde{x} \cdot \tilde{\beta} = 3\beta_1 + 5\beta_2$, and by Theorem 0.2:

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\tilde{x} \cdot \tilde{\beta})}_{2 \times 1} = \underbrace{\tilde{x}}_{2 \times 1} = \begin{pmatrix} 3 \\ 5 \end{pmatrix}$$

Verifying entry-wise:

$$\begin{aligned} \frac{\partial}{\partial \tilde{\beta}}(3\beta_1 + 5\beta_2) &= \begin{pmatrix} \frac{\partial}{\partial \beta_1}(3\beta_1 + 5\beta_2) \\ \frac{\partial}{\partial \beta_2}(3\beta_1 + 5\beta_2) \end{pmatrix} \\ &= \begin{pmatrix} 3 \\ 5 \end{pmatrix} \end{aligned}$$

Both methods agree.

Theorem 0.3 (Product rule for dot-products). *If $\tilde{a} = \tilde{a}(\tilde{x})$ and $\tilde{b} = \tilde{b}(\tilde{x})$ are differentiable $p \times 1$ vector functions of $\tilde{x}$, then:*

$$\frac{\partial}{\partial \tilde{x}} \underbrace{\tilde{x}}_{p \times 1} \cdot \underbrace{\tilde{b}}_{p \times 1} = \left(\frac{\partial}{\partial \tilde{x}} \underbrace{\tilde{x}}_{p \times 1} \underbrace{\tilde{a}^\top}_{1 \times p} \right) \underbrace{\tilde{b}}_{p \times 1} + \left(\frac{\partial}{\partial \tilde{x}} \underbrace{\tilde{x}}_{p \times 1} \underbrace{\tilde{b}^\top}_{1 \times p} \right) \underbrace{\tilde{a}}_{p \times 1}$$

i Proof

Proof. Entry-wise, for $i = 1, \dots, p$:

$$\begin{aligned} \left[\frac{\partial}{\partial \tilde{x}}(\tilde{a} \cdot \tilde{b}) \right]_i &= \frac{\partial}{\partial x_i} \sum_{k=1}^p a_k b_k \\ &= \sum_{k=1}^p \left(b_k \frac{\partial}{\partial x_i} a_k + a_k \frac{\partial}{\partial x_i} b_k \right) \\ &= \left[\left(\frac{\partial}{\partial \tilde{x}} \tilde{a}^\top \right) \tilde{b} \right]_i + \left[\left(\frac{\partial}{\partial \tilde{x}} \tilde{b}^\top \right) \tilde{a} \right]_i \end{aligned}$$

□

Example 0.11 (Example of the dot-product rule). Apply Theorem 0.3 with the vector $\tilde{\beta} = (\beta_1, \beta_2)^\top$ in the role of $\tilde{x}$. Let $\tilde{a}(\tilde{\beta}) = (\beta_1, \beta_1\beta_2)^\top$ and $\tilde{b}(\tilde{\beta}) = (\beta_2, \beta_1)^\top$. Then:

$$\begin{aligned} \tilde{a} \cdot \tilde{b} &= \beta_1 \cdot \beta_2 + \beta_1\beta_2 \cdot \beta_1 \\ &= \beta_1\beta_2 + \beta_1^2\beta_2 \end{aligned}$$

By direct calculation:

$$\begin{aligned}\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\tilde{a} \cdot \tilde{b})}_{2 \times 1} &= \frac{\partial}{\partial \tilde{\beta}}(\beta_1 \beta_2 + \beta_1^2 \beta_2) \\ &= \begin{pmatrix} \beta_2 + 2\beta_1 \beta_2 \\ \beta_1 + \beta_1^2 \end{pmatrix}\end{aligned}$$

By the product rule (Theorem 0.3), using $\underbrace{\frac{\partial}{\partial \tilde{\beta}} \tilde{a}^\top}_{2 \times 2} = \begin{pmatrix} 1 & \beta_2 \\ 0 & \beta_1 \end{pmatrix}$ and $\underbrace{\frac{\partial}{\partial \tilde{\beta}} \tilde{b}^\top}_{2 \times 2} = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$:

$$\begin{aligned}\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\tilde{a} \cdot \tilde{b})}_{2 \times 1} &= \underbrace{\begin{pmatrix} 1 & \beta_2 \\ 0 & \beta_1 \end{pmatrix}}_{2 \times 2} \underbrace{\begin{pmatrix} \beta_2 \\ \beta_1 \end{pmatrix}}_{2 \times 1} + \underbrace{\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}}_{2 \times 2} \underbrace{\begin{pmatrix} \beta_1 \\ \beta_1 \beta_2 \end{pmatrix}}_{2 \times 1} \\ &= \begin{pmatrix} \beta_2 + \beta_1 \beta_2 \\ \beta_1^2 \end{pmatrix} + \begin{pmatrix} \beta_1 \beta_2 \\ \beta_1 \end{pmatrix} \\ &= \begin{pmatrix} \beta_2 + 2\beta_1 \beta_2 \\ \beta_1^2 + \beta_1 \end{pmatrix}\end{aligned}$$

Both methods agree.

Theorem 0.4 (Derivative of a linear map). *If $\mathbf{A}$ is an $m \times p$ matrix that is constant with respect to $\tilde{\beta}$, then:*

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\mathbf{A} \tilde{\beta})}_{p \times m} = \underbrace{\mathbf{A}^\top}_{p \times m}$$

i Proof

Proof. For entry (i, j) , where row i indexes the denominator $\tilde{\beta}$ (see Definition 0.6) and column j indexes the numerator $\mathbf{A} \tilde{\beta}$:

$$\begin{aligned}\left[\frac{\partial}{\partial \tilde{\beta}}(\mathbf{A} \tilde{\beta}) \right]_{ij} &= \frac{\partial}{\partial \beta_i}(\mathbf{A} \tilde{\beta})_j \\ &= \frac{\partial}{\partial \beta_i} \sum_{k=1}^p a_{jk} \beta_k \\ &= a_{ji} \\ &= [\mathbf{A}^\top]_{ij}\end{aligned}$$

□

Example 0.12 (Derivative of a linear map). Let $\mathbf{A} = \begin{pmatrix} 2 & 3 \end{pmatrix}$ (1×2) and $\tilde{\beta} = (\beta_1, \beta_2)^\top$. Then $\mathbf{A}\tilde{\beta} = 2\beta_1 + 3\beta_2$, and by Theorem 0.4:

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\mathbf{A}\tilde{\beta})}_{2 \times 1} = \underbrace{\mathbf{A}^\top}_{2 \times 1} = \begin{pmatrix} 2 \\ 3 \end{pmatrix}$$

Theorem 0.5 (Vector-derivative of a matrix-vector product). *If $\mathbf{A}$ is an $m \times q$ matrix that is constant with respect to $\tilde{\beta}$, and $\tilde{v} = \tilde{v}(\tilde{\beta})$ is a $q \times 1$ vector that depends on the $p \times 1$ vector $\tilde{\beta}$, then:*

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\mathbf{A}\tilde{v})}_{p \times m} = \underbrace{\left(\frac{\partial}{\partial \tilde{\beta}} \tilde{v} \right)}_{p \times q} \underbrace{\mathbf{A}^\top}_{q \times m}$$

i Proof

Proof. For entry (i, j) , where row i indexes the denominator $\tilde{\beta}$ and column j indexes the numerator $\mathbf{A}\tilde{v}$ (see Definition 0.6):

$$\begin{aligned} \left[\frac{\partial}{\partial \tilde{\beta}}(\mathbf{A}\tilde{v}) \right]_{ij} &= \frac{\partial}{\partial \beta_i}(\mathbf{A}\tilde{v})_j \\ &= \frac{\partial}{\partial \beta_i} \sum_{k=1}^q a_{jk} v_k \\ &= \sum_{k=1}^q a_{jk} \frac{\partial}{\partial \beta_i} v_k \\ &= \sum_{k=1}^q \left[\frac{\partial}{\partial \tilde{\beta}} \tilde{v} \right]_{ik} [\mathbf{A}^\top]_{kj} \\ &= \left[\left(\frac{\partial}{\partial \tilde{\beta}} \tilde{v} \right) \mathbf{A}^\top \right]_{ij} \end{aligned}$$

□

Example 0.13 (Vector-derivative of a matrix-vector product). Let $\mathbf{A} = \begin{pmatrix} 2 & 3 \end{pmatrix}$ (1×2 , constant) and $\tilde{v}(\tilde{\beta}) = (\beta_1^2, \beta_2^2)^\top$. Then $\mathbf{A}\tilde{v} = 2\beta_1^2 + 3\beta_2^2$. By Theorem 0.5:

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\mathbf{A}\tilde{v})}_{2 \times 1} = \begin{pmatrix} 2\beta_1 & 0 \\ 0 & 2\beta_2 \end{pmatrix} \begin{pmatrix} 2 \\ 3 \end{pmatrix} = \begin{pmatrix} 4\beta_1 \\ 6\beta_2 \end{pmatrix}$$

Remark 0.3 (The derivative of a linear map as a special case). This result generalizes Theorem 0.4, which is the special case $\tilde{v} = \tilde{\beta}$ (so that $q = p$, $\frac{\partial}{\partial \tilde{\beta}}\tilde{\beta} = \mathbf{I}_p$, and $\frac{\partial}{\partial \tilde{\beta}}(\mathbf{A}\tilde{\beta}) = \mathbf{I}_p \mathbf{A}^\top$, which is $\mathbf{A}^\top$). For example, Example 0.12 is the case $\mathbf{A} = \begin{pmatrix} 2 & 3 \end{pmatrix}$ and

$$\begin{aligned} \tilde{v} &= \tilde{\beta} \\ &= (\beta_1, \beta_2)^\top, \end{aligned}$$

where this result gives $\mathbf{I}_2 \mathbf{A}^\top = (2, 3)^\top$.

Theorem 0.6 (Vector-derivative of a product of matrices). *If $\mathbf{A}$ ($k \times m$) and $\mathbf{B}$ ($m \times q$) are constant with respect to $\tilde{\beta}$, and $\tilde{v} = \tilde{v}(\tilde{\beta})$ is a $q \times 1$ vector that depends on the $p \times 1$ vector $\tilde{\beta}$, then:*

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\mathbf{A}\mathbf{B}\tilde{v})}_{p \times k} = \underbrace{\left(\frac{\partial}{\partial \tilde{\beta}} \tilde{v} \right)}_{p \times q} \underbrace{\mathbf{B}^\top}_{q \times m} \underbrace{\mathbf{A}^\top}_{m \times k}$$

i Proof

Proof. Apply Theorem 0.5 with the constant $k \times q$ matrix $\mathbf{A}\mathbf{B}$, then use the [transpose of a product](#):

$$\begin{aligned} \frac{\partial}{\partial \tilde{\beta}}(\mathbf{A}\mathbf{B}\tilde{v}) &= \left(\frac{\partial}{\partial \tilde{\beta}} \tilde{v} \right) (\mathbf{A}\mathbf{B})^\top \quad (\text{derivative of a matrix-vector product}) \\ &= \left(\frac{\partial}{\partial \tilde{\beta}} \tilde{v} \right) \mathbf{B}^\top \mathbf{A}^\top \quad (\text{transpose of a product}) \end{aligned}$$

□

Example 0.14. Let $\mathbf{A} = \begin{pmatrix} 1 & 0 \end{pmatrix}$ (1×2), $\mathbf{B} = \begin{pmatrix} 2 & 0 \\ 0 & 3 \end{pmatrix}$ (2×2), and $\tilde{v}(\tilde{\beta}) = \tilde{\beta}$ where $\tilde{\beta} = (\beta_1, \beta_2)^\top$. Then $\mathbf{AB}\tilde{v} = 2\beta_1$, and:

$$\begin{aligned} \underbrace{\frac{\partial}{\partial \tilde{\beta}}(\mathbf{AB}\tilde{v})}_{2 \times 1} &= \underbrace{\left(\frac{\partial}{\partial \tilde{\beta}} \tilde{\beta} \right)}_{2 \times 2} \underbrace{\mathbf{B}^\top}_{2 \times 2} \underbrace{\mathbf{A}^\top}_{2 \times 1} \\ &= \mathbf{I}_2 \begin{pmatrix} 2 & 0 \\ 0 & 3 \end{pmatrix} \begin{pmatrix} 1 \\ 0 \end{pmatrix} \\ &= \begin{pmatrix} 2 \\ 0 \end{pmatrix} \end{aligned}$$

Corollary 0.1 (Derivative of a dot product, transpose-product form). *If $\tilde{x}$ is constant with respect to $\tilde{\beta}$, then:*

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\underbrace{\tilde{x}^\top}_{1 \times p} \underbrace{\tilde{\beta}}_{p \times 1})}_{p \times 1} = \underbrace{\tilde{x}}_{p \times 1}$$

i Proof

Proof. Using Theorem 0.2:

Since $\tilde{x}^\top \tilde{\beta} = \tilde{x} \cdot \tilde{\beta}$ (see [dot product](#)), and $\tilde{x}$ is constant with respect to $\tilde{\beta}$:

$$\begin{aligned} \frac{\partial}{\partial \tilde{\beta}}(\tilde{x}^\top \tilde{\beta}) &= \frac{\partial}{\partial \tilde{\beta}}(\tilde{x} \cdot \tilde{\beta}) \\ &= \tilde{x} \end{aligned}$$

by Theorem 0.2. □

i Proof

Proof. Using Theorem 0.5:

Since $\tilde{x}$ is constant with respect to $\tilde{\beta}$, $\mathbf{A} = \tilde{x}^\top$ is a constant $1 \times p$ matrix. Applying Theorem 0.5 with $\tilde{v} = \tilde{\beta}$ (so $\frac{\partial}{\partial \tilde{\beta}} \tilde{\beta} = \mathbf{I}$):

$$\begin{aligned}
\frac{\partial}{\partial \tilde{\beta}}(\tilde{x}^\top \tilde{\beta}) &= \left(\frac{\partial}{\partial \tilde{\beta}} \tilde{\beta} \right) (\tilde{x}^\top)^\top \\
&= \mathbf{I} \cdot \tilde{x} \\
&= \tilde{x}
\end{aligned}$$

□

Example 0.15 (Derivative of a transpose product). Let $\tilde{x} = (3, 5)^\top$ and $\tilde{\beta} = (\beta_1, \beta_2)^\top$. Then $\tilde{x}^\top \tilde{\beta} = 3\beta_1 + 5\beta_2$, and by Corollary 0.1:

$$\begin{aligned}
\underbrace{\frac{\partial}{\partial \tilde{\beta}} \left(\underbrace{\tilde{x}^\top}_{1 \times 2} \underbrace{\tilde{\beta}}_{2 \times 1} \right)}_{2 \times 1} &= \underbrace{\tilde{x}}_{2 \times 1} \\
&= \begin{pmatrix} 3 \\ 5 \end{pmatrix}
\end{aligned}$$

Remark 0.4 (The coefficient gets transposed). This vector derivative formula looks a lot like non-vector calculus, except that you have to transpose the coefficient: in scalar calculus $\frac{\partial}{\partial x}(cx) = c$, but here the coefficient $\tilde{x}^\top$ (a row vector) becomes $\tilde{x}$ (a column vector) in the result. For example, with $\tilde{x} = (2, -1)^\top$, $\tilde{x}^\top \tilde{\beta} = 2\beta_1 - \beta_2$, whose vector derivative is the column vector $(2, -1)^\top = \tilde{x}$, not the row vector $\tilde{x}^\top = (2, -1)$.

Theorem 0.7 (Derivative of a quadratic form). *For a quadratic form (see [quadratic form](#)), if $\mathbf{S}$ is a symmetric $p \times p$ matrix that is constant with respect to $\tilde{\beta}$, then:*

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\tilde{\beta}^\top \mathbf{S} \tilde{\beta})}_{p \times 1} = \underbrace{2\mathbf{S}\tilde{\beta}}_{p \times 1}$$

i Proof

Proof. Expanding entry-wise, $\tilde{\beta}^\top \mathbf{S} \tilde{\beta} = \sum_{j=1}^p \sum_{k=1}^p s_{jk} \beta_j \beta_k$. Differentiating component-wise with respect to β_i for $i = 1, \dots, p$:

$$\begin{aligned}
\left[\frac{\partial}{\partial \tilde{\beta}} (\tilde{\beta}^\top \mathbf{S} \tilde{\beta}) \right]_i &= \frac{\partial}{\partial \beta_i} \sum_{j=1}^p \sum_{k=1}^p s_{jk} \beta_j \beta_k \quad (\text{expand quadratic form}) \\
&= \sum_{k=1}^p s_{ik} \beta_k + \sum_{j=1}^p s_{ji} \beta_j \quad (\text{product rule for } \beta_i \beta_k) \\
&= [\mathbf{S} \tilde{\beta}]_i + [\mathbf{S}^\top \tilde{\beta}]_i \quad (\text{matrix-vector multiplication definition}) \\
&= [(\mathbf{S} + \mathbf{S}^\top) \tilde{\beta}]_i \quad (\text{linearity of matrix multiplication})
\end{aligned}$$

When $\mathbf{S}$ is symmetric ($\mathbf{S} = \mathbf{S}^\top$), $\mathbf{S} + \mathbf{S}^\top = 2\mathbf{S}$, so $\frac{\partial}{\partial \tilde{\beta}} (\tilde{\beta}^\top \mathbf{S} \tilde{\beta}) = 2\mathbf{S} \tilde{\beta}$. $\square$

Remark 0.5 (Like the derivative of cx^2). This operation is like taking the derivative of cx^2 with respect to x in non-vector calculus: $\frac{\partial}{\partial x}(cx^2) = 2cx$, and Theorem 0.7 says $\frac{\partial}{\partial \tilde{\beta}}(\tilde{\beta}^\top \mathbf{S} \tilde{\beta}) = 2\mathbf{S} \tilde{\beta}$. For example, with $p = 1$, $\mathbf{S} = (3)$, and $\tilde{\beta} = (\beta_1)$, $\tilde{\beta}^\top \mathbf{S} \tilde{\beta} = 3\beta_1^2$, and its derivative is $6\beta_1 = 2\mathbf{S} \tilde{\beta}$.

Example 0.16 (Derivative of a quadratic form). Let $\mathbf{S} = \begin{pmatrix} 3 & 1 \\ 1 & 2 \end{pmatrix}$ (2×2 , symmetric and constant) and $\tilde{\beta} = (\beta_1, \beta_2)^\top$. Then $\tilde{\beta}^\top \mathbf{S} \tilde{\beta} = 3\beta_1^2 + 2\beta_1\beta_2 + 2\beta_2^2$. By Theorem 0.7:

$$\begin{aligned}
\underbrace{\frac{\partial}{\partial \tilde{\beta}} (\tilde{\beta}^\top \mathbf{S} \tilde{\beta})}_{2 \times 1} &= 2\mathbf{S} \tilde{\beta} \\
&= 2 \begin{pmatrix} 3 & 1 \\ 1 & 2 \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \end{pmatrix} \\
&= \begin{pmatrix} 6\beta_1 + 2\beta_2 \\ 2\beta_1 + 4\beta_2 \end{pmatrix}
\end{aligned}$$

Differentiating component-wise directly:

$$\begin{pmatrix} \frac{\partial}{\partial \beta_1} (3\beta_1^2 + 2\beta_1\beta_2 + 2\beta_2^2) \\ \frac{\partial}{\partial \beta_2} (3\beta_1^2 + 2\beta_1\beta_2 + 2\beta_2^2) \end{pmatrix} = \begin{pmatrix} 6\beta_1 + 2\beta_2 \\ 2\beta_1 + 4\beta_2 \end{pmatrix}$$

Both methods agree.

Corollary 0.2 (Derivative of a simple quadratic form).

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}}_{p \times 1}(\tilde{\beta} \cdot \tilde{\beta}) = \underbrace{2\tilde{\beta}}_{p \times 1}$$

i Proof

Proof. Applying Theorem 0.7 with $\mathbf{S} = \mathbf{I}_{p \times p}$ (which is symmetric and constant with respect to $\tilde{\beta}$):

$$\begin{aligned} \frac{\partial}{\partial \tilde{\beta}}(\tilde{\beta} \cdot \tilde{\beta}) &= \frac{\partial}{\partial \tilde{\beta}}(\tilde{\beta}^\top \mathbf{I}_{p \times p} \tilde{\beta}) \quad (\text{rewrite with identity matrix}) \\ &= 2\mathbf{I}_{p \times p} \tilde{\beta} \quad (\text{derivative of a quadratic form, with } \mathbf{S} = \mathbf{I}_{p \times p}) \\ &= 2\tilde{\beta} \quad (\text{identity matrix property}) \end{aligned}$$

□

Remark 0.6 (Like the derivative of x^2). This vector derivative is like taking the derivative of x^2 : in scalar calculus $\frac{\partial}{\partial x} x^2 = 2x$, and Corollary 0.2 says $\frac{\partial}{\partial \tilde{\beta}}(\tilde{\beta} \cdot \tilde{\beta}) = 2\tilde{\beta}$. For example, with $p = 1$ and $\tilde{\beta} = (\beta_1)$, $\tilde{\beta} \cdot \tilde{\beta} = \beta_1^2$, and its derivative is $2\beta_1 = 2\tilde{\beta}$.

Example 0.17 (Derivative of a sum of squares). Let $\tilde{\beta} = (\beta_1, \beta_2)^\top$, so $\tilde{\beta} \cdot \tilde{\beta} = \beta_1^2 + \beta_2^2$. By Corollary 0.2:

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}}_{2 \times 1}(\tilde{\beta} \cdot \tilde{\beta}) = 2\tilde{\beta} = \begin{pmatrix} 2\beta_1 \\ 2\beta_2 \end{pmatrix}$$

Direct partial differentiation yields the same column vector.

Definition 0.9 (Residuals and squared errors). Let $\mathbf{X}$ be an $n \times p$ matrix, $\tilde{y} \in \mathbb{R}^n$, and $\tilde{\beta} \in \mathbb{R}^p$. The **residual vector** of $\tilde{\beta}$ is $\tilde{\varepsilon}(\tilde{\beta}) \stackrel{\text{def}}{=} \tilde{y} - \mathbf{X}\tilde{\beta}$, and its entries $\varepsilon_i = y_i - (\mathbf{X}\tilde{\beta})_i$ are the **residuals**. The square ε_i^2 is the i th **squared error**, and their sum

$$\begin{aligned}\tilde{\varepsilon} \cdot \tilde{\varepsilon} &= \sum_{i=1}^n \tilde{\varepsilon}_i^2 \\ &= \|\tilde{\mathbf{y}} - \mathbf{X}\tilde{\boldsymbol{\beta}}\|^2\end{aligned}$$

is the **residual sum of squares**. A **least squares solution** of $\mathbf{X}\tilde{\boldsymbol{\beta}} = \tilde{\mathbf{y}}$ makes $\|\tilde{\mathbf{y}} - \mathbf{X}\tilde{\boldsymbol{\beta}}\|$ as small as possible, and so makes the residual sum of squares as small as possible too.

Example 0.18 (Residuals of a constant fit). Let $\mathbf{X} = (1, 1)^\top$ ($n = 2$, $p = 1$) and $\tilde{\mathbf{y}} = (1, 3)^\top$. For $\tilde{\boldsymbol{\beta}} = (2)$, $\mathbf{X}\tilde{\boldsymbol{\beta}} = (2, 2)^\top$, so

$$\begin{aligned}\tilde{\varepsilon}(\tilde{\boldsymbol{\beta}}) &= (1, 3)^\top - (2, 2)^\top \quad (\text{definition of the residual vector}) \\ &= (-1, 1)^\top, \quad (\text{subtract entry by entry})\end{aligned}$$

the squared errors are $(-1)^2 = 1$ and $1^2 = 1$, and the residual sum of squares is $1 + 1 = 2$. For $\tilde{\boldsymbol{\beta}} = (1)$ the residuals are 0 and 2, and the residual sum of squares is $0 + 4 = 4$.

Videos

Jon Krohn’s “Calculus for Machine Learning” YouTube playlist has videos on gradients of sums and means of squared errors (Definition 0.9):

- [The Gradient of Quadratic Cost](#)
- [The Gradient of Mean Squared Error](#)

Theorem 0.8 (Vector chain rule). *Let $\tilde{\mathbf{x}}$ be a $p \times 1$ vector, let $\tilde{\mathbf{y}} = \tilde{\mathbf{g}}(\tilde{\mathbf{x}})$ be a $q \times 1$ vector-valued function of $\tilde{\mathbf{x}}$, and let $z = f(\tilde{\mathbf{y}})$ be a scalar-valued function of $\tilde{\mathbf{y}}$, where $\tilde{\mathbf{g}}$ and f have continuous partial derivatives. Then $z = f(\tilde{\mathbf{g}}(\tilde{\mathbf{x}}))$, as a function of $\tilde{\mathbf{x}}$, satisfies*

$$\frac{\partial z}{\partial \tilde{\mathbf{x}}} = \frac{\partial \tilde{\mathbf{y}}}{\partial \tilde{\mathbf{x}}} \frac{\partial z}{\partial \tilde{\mathbf{y}}}$$

$p \times 1 \quad p \times q \quad q \times 1$

where $\frac{\partial \tilde{\mathbf{y}}}{\partial \tilde{\mathbf{x}}}$ is the derivative of Definition 0.6 and $\frac{\partial z}{\partial \tilde{\mathbf{x}}}$ and $\frac{\partial z}{\partial \tilde{\mathbf{y}}}$ are vector derivatives (Definition 0.2).

Remark 0.7 (The order of the factors matters). The vector chain rule (Theorem 0.8) is like the univariate **chain rule**, but the order matters now: $\frac{\partial \tilde{\mathbf{y}}}{\partial \tilde{\mathbf{x}}}$ is $p \times q$ and $\frac{\partial z}{\partial \tilde{\mathbf{y}}}$ is $q \times 1$, so the product $\frac{\partial z}{\partial \tilde{\mathbf{y}}} \frac{\partial \tilde{\mathbf{y}}}{\partial \tilde{\mathbf{x}}}$ in the other order is not even defined unless $p = 1$. The version presented here is for the gradient (Definition 0.2), a column vector. The total

derivative (Definition 0.5), a row vector, is the transpose of the gradient (Theorem 0.1), and transposing both sides gives $\left(\frac{\partial z}{\partial \tilde{x}}\right)^\top = \left(\frac{\partial z}{\partial \tilde{y}}\right)^\top \left(\frac{\partial \tilde{y}}{\partial \tilde{x}}\right)^\top$, with the factors in the reverse order; there $\left(\frac{\partial \tilde{y}}{\partial \tilde{x}}\right)^\top$ is the Jacobian matrix of $\tilde{y}$ as a function of $\tilde{x}$ (Definition 0.7).

Example 0.19 (Applying the vector chain rule). Let $\tilde{x} = (x_1, x_2)^\top$,

$$\begin{aligned}\tilde{y} &= \tilde{g}(\tilde{x}) \\ &= (x_1 + x_2, x_1 x_2)^\top,\end{aligned}$$

and

$$\begin{aligned}z &= f(\tilde{y}) \\ &= y_1^2 + y_2.\end{aligned}$$

Then

$$\begin{aligned}\underbrace{\frac{\partial z}{\partial \tilde{x}}}_{2 \times 1} &= \underbrace{\frac{\partial \tilde{y}}{\partial \tilde{x}}}_{2 \times 2} \underbrace{\frac{\partial z}{\partial \tilde{y}}}_{2 \times 1} && \text{(vector chain rule)} \\ &= \begin{bmatrix} 1 & x_2 \\ 1 & x_1 \end{bmatrix} \begin{bmatrix} 2y_1 \\ 1 \end{bmatrix} && \text{(differentiate } \tilde{g} \text{ and } f) \\ &= \begin{bmatrix} 2(x_1 + x_2) + x_2 \\ 2(x_1 + x_2) + x_1 \end{bmatrix} && \text{(multiply, and substitute } y_1 = x_1 + x_2) \\ &= \begin{bmatrix} 2x_1 + 3x_2 \\ 3x_1 + 2x_2 \end{bmatrix} && \text{(collect terms)}\end{aligned}$$

This matches differentiating z directly:

$$\begin{aligned}z &= (x_1 + x_2)^2 + x_1 x_2 \\ &= x_1^2 + 3x_1 x_2 + x_2^2,\end{aligned}$$

so $\frac{\partial}{\partial x_1} z = 2x_1 + 3x_2$ and $\frac{\partial}{\partial x_2} z = 3x_1 + 2x_2$. The product in the other order, a 2×1 matrix times a 2×2 matrix, is not defined. The total derivative is the transpose, the row vector $(2x_1 + 3x_2, 3x_1 + 2x_2)$.

See Felippa (n.d.), and Wikipedia contributors (n.d.), section “Relationship with Frechet derivative”.

Corollary 0.3 (Vector chain rule for quadratic forms). *If $\tilde{\varepsilon} = \tilde{\varepsilon}(\tilde{\beta})$ is an $n \times 1$ vector-valued function of the $p \times 1$ vector $\tilde{\beta}$, with continuous partial derivatives, then*

$$\underbrace{\frac{\partial}{\partial \tilde{\beta}}(\tilde{\varepsilon}(\tilde{\beta}) \cdot \tilde{\varepsilon}(\tilde{\beta}))}_{p \times 1} = \underbrace{\left(\frac{\partial}{\partial \tilde{\beta}} \tilde{\varepsilon}(\tilde{\beta}) \right)}_{p \times n} \underbrace{(2\tilde{\varepsilon}(\tilde{\beta}))}_{n \times 1}$$

Proof

Proof. Apply Theorem 0.8 with $\tilde{x} = \tilde{\beta}$, $\tilde{y} = \tilde{\varepsilon}$, $q = n$, and

$$\begin{aligned} z &= f(\tilde{\varepsilon}) \\ &= \tilde{\varepsilon} \cdot \tilde{\varepsilon} \\ &= \tilde{\varepsilon}^\top \tilde{\varepsilon} : \end{aligned}$$

$$\begin{aligned} \frac{\partial}{\partial \tilde{\beta}}(\tilde{\varepsilon} \cdot \tilde{\varepsilon}) &= \left(\frac{\partial}{\partial \tilde{\beta}} \tilde{\varepsilon} \right) \frac{\partial}{\partial \tilde{\varepsilon}}(\tilde{\varepsilon}^\top \tilde{\varepsilon}) \quad (\text{vector chain rule}) \\ &= \left(\frac{\partial}{\partial \tilde{\beta}} \tilde{\varepsilon} \right) (2\tilde{\varepsilon}) \quad (\text{derivative of a simple quadratic form, in } \tilde{\varepsilon}) \end{aligned}$$

The second step is Corollary 0.2, applied to the $n \times 1$ vector $\tilde{\varepsilon}$ in place of $\tilde{\beta}$. □

Example 0.20 (Derivative of the residual sum of squares). Let $\tilde{y}$ ($n \times 1$) and $\mathbf{X}$ ($n \times p$) be constant with respect to $\tilde{\beta}$, and let $\tilde{\varepsilon}(\tilde{\beta}) = \tilde{y} - \mathbf{X}\tilde{\beta}$ be the residual vector (Definition 0.9). By Theorem 0.4, $\frac{\partial}{\partial \tilde{\beta}}(\mathbf{X}\tilde{\beta}) = \mathbf{X}^\top$, and $\frac{\partial}{\partial \tilde{\beta}}\tilde{y} = \mathbf{0}_{p \times n}$ because $\tilde{y}$ is constant, so $\frac{\partial}{\partial \tilde{\beta}}\tilde{\varepsilon} = -\mathbf{X}^\top$ ($p \times n$). By Corollary 0.3:

$$\begin{aligned} \frac{\partial}{\partial \tilde{\beta}}(\tilde{\varepsilon} \cdot \tilde{\varepsilon}) &= (-\mathbf{X}^\top)(2\tilde{\varepsilon}) \quad (\text{vector chain rule for quadratic forms}) \\ &= -2\mathbf{X}^\top(\tilde{y} - \mathbf{X}\tilde{\beta}) \quad (\text{substitute } \tilde{\varepsilon} = \tilde{y} - \mathbf{X}\tilde{\beta}) \end{aligned}$$

Setting this $p \times 1$ vector to $\tilde{0}$ gives the normal equations $\mathbf{X}^\top \mathbf{X}\tilde{\beta} = \mathbf{X}^\top \tilde{y}$ of least squares.

Videos

Jon Krohn’s “Calculus for Machine Learning” YouTube playlist has videos on the chain rule for partial derivatives:

- [The Chain Rule for Partial Derivatives](#)
- [Exercises on the Multivariate Chain Rule](#)

Definition 0.10 (Matrix derivative (gradient)). For a scalar-valued function $f(\mathbf{X})$ of an $m \times n$ matrix $\mathbf{X}$, the **matrix derivative** is the $m \times n$ matrix whose (i, j) entry is the partial derivative of f with respect to the (i, j) entry of $\mathbf{X}$:

$$\left[\frac{\partial}{\partial \mathbf{X}} f \right]_{ij} = \frac{\partial}{\partial X_{ij}} f$$

Like the vector derivative (Definition 0.2), the matrix derivative is also called the **gradient** of f with respect to $\mathbf{X}$, written $\nabla_{\mathbf{X}} f$.

Example 0.21 (The matrix derivative of a trace). Let $\mathbf{X}$ be a 2×2 matrix and

$$\begin{aligned} f(\mathbf{X}) &= \text{tr}(\mathbf{X}) \\ &= X_{11} + X_{22} \end{aligned}$$

(see [trace](#)). Then $\frac{\partial}{\partial X_{ij}} f = 1$ if $i = j$ and 0 otherwise, so:

$$\frac{\partial}{\partial \mathbf{X}} f = \mathbf{I}_2$$

Theorem 0.9 (Matrix derivative of the trace of a matrix product). If $\mathbf{A}$ ($r \times m$) and $\mathbf{B}$ ($n \times r$) are constant with respect to the $m \times n$ matrix $\mathbf{X}$, then:

$$\underbrace{\frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{A}\mathbf{X}\mathbf{B})}_{m \times n} = \underbrace{\mathbf{A}^\top}_{m \times r} \underbrace{\mathbf{B}^\top}_{r \times n}$$

i Proof

Proof. Write A_{kl} , X_{kl} , and B_{kl} for the entries of $\mathbf{A}$, $\mathbf{X}$, and $\mathbf{B}$. For entry (i, j) :

$$\begin{aligned} \left[\frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{A}\mathbf{X}\mathbf{B}) \right]_{ij} &= \frac{\partial}{\partial X_{ij}} \sum_{a=1}^r \sum_{b=1}^m \sum_{c=1}^n A_{ab} X_{bc} B_{ca} \quad (\text{trace of the } r \times r \text{ product } \mathbf{A}\mathbf{X}\mathbf{B}) \\ &= \sum_{a=1}^r A_{ai} B_{ja} \quad (\text{only the terms with } b = i, c = j \text{ depend on } X_{ij}) \\ &= \sum_{a=1}^r [\mathbf{A}^\top]_{ia} [\mathbf{B}^\top]_{aj} \quad (\text{definition of the transpose}) \\ &= [\mathbf{A}^\top \mathbf{B}^\top]_{ij} \quad (\text{definition of matrix multiplication}) \end{aligned}$$

□

Remark 0.8 (Why the theorem uses the trace). The trace makes $\text{tr}(\mathbf{AXB})$ a scalar, so its matrix derivative (Definition 0.10) is again an $m \times n$ matrix. The matrix product $\mathbf{AXB}$ itself (without the trace) is an $r \times r$ matrix, and each of its r^2 entries has a partial derivative with respect to each of the mn entries of $\mathbf{X}$. Those r^2mn partial derivatives form a four-index array (a fourth-order tensor), not a matrix, which is why this result is stated for the scalar $\text{tr}(\mathbf{AXB})$.

For example, with

$$\begin{aligned}\mathbf{A} &= \mathbf{B} \\ &= \mathbf{I}_2,\end{aligned}$$

the product $\mathbf{AXB} = \mathbf{X}$ has 4 entries, each with 4 partial derivatives, 16 in all, while its trace $X_{11} + X_{22}$ has the 4 partial derivatives that form the 2×2 matrix $\mathbf{I}_2$ of Example 0.21.

Example 0.22 (Differentiating a weighted trace). Let $\mathbf{A} = \mathbf{I}_2$ (2×2) and $\mathbf{B} = \begin{pmatrix} 2 & 0 \\ 0 & 3 \end{pmatrix}$ (2×2). Then $\text{tr}(\mathbf{AXB}) = 2X_{11} + 3X_{22}$, and:

$$\begin{aligned}\underbrace{\frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{AXB})}_{2 \times 2} &= \underbrace{\mathbf{A}^\top}_{2 \times 2} \underbrace{\mathbf{B}^\top}_{2 \times 2} \\ &= \mathbf{I}_2 \begin{pmatrix} 2 & 0 \\ 0 & 3 \end{pmatrix} \\ &= \begin{pmatrix} 2 & 0 \\ 0 & 3 \end{pmatrix}\end{aligned}$$

Exercise 0.7 (A gradient that does not mention its variable). Fix

$$A = \begin{bmatrix} 2 & -1 & 0 \\ 4 & 3 & -2 \end{bmatrix}$$

let $W \in \mathbb{R}^{2 \times 3}$ vary, and define

$$g(W) = \sum_{i=1}^2 \sum_{j=1}^3 W_{ij} A_{ij}$$

which is the entry-by-entry version of the inner product for matrices. Compute $\nabla_W g(W)$, and say what is unusual about the answer.

Solution

Solution 0.7. Written out over the six positions,

$$g(W) = 2W_{11} - W_{12} + 0W_{13} + 4W_{21} + 3W_{22} - 2W_{23}$$

Each W_{ij} appears in exactly one term, multiplied by A_{ij} and by nothing else, so differentiating with respect to it leaves A_{ij} behind:

$$\frac{\partial g}{\partial W_{ij}} = A_{ij}$$

Collecting those partial derivatives into a matrix of the same shape as W , as Table 1 requires,

$$\begin{aligned}\nabla_W g(W) &= \begin{bmatrix} 2 & -1 & 0 \\ 4 & 3 & -2 \end{bmatrix} \\ &= A\end{aligned}$$

What is unusual is that the gradient came out **constant**: it does not mention W at all. The constant answer is not a quirk of this particular A . g is a linear function of W , and the gradient of a linear function is constant everywhere, for the same reason the derivative of $f(w) = cw$ is c no matter where it is evaluated.

A constant gradient is the easy case, and it is not the case we usually face. Most functions we minimize are curved — the residual sum of squares of a linear model (Definition 0.9) is the standard example — so their gradient changes from point to point and the downhill direction has to be worked out afresh at every step.

Second partial derivatives and the Hessian

Definition 0.11 (Open and closed balls). Let $\tilde{x} \in \mathbb{R}^p$ and $r > 0$. The **open ball** of radius r around $\tilde{x}$ is the set $\{\tilde{y} \in \mathbb{R}^p : \|\tilde{y} - \tilde{x}\| < r\}$, of points whose distance from $\tilde{x}$ (the **Euclidean norm** of the difference) is less than r . The **closed ball** of radius r around $\tilde{x}$ is $\{\tilde{y} \in \mathbb{R}^p : \|\tilde{y} - \tilde{x}\| \leq r\}$. The open and closed balls of radius 1 around $\tilde{0}$ are the open and closed **unit balls**.

Example 0.23 (Points in the unit balls of $\mathbb{R}^2$).

- $(0.5, 0.5)^\top$ has

$$\begin{aligned}\|(0.5, 0.5)^\top\|^2 &= 0.25 + 0.25 \\ &= 0.5 \\ &< 1,\end{aligned}$$

so it is in both the open and the closed unit ball.

- $(0.6, 0.8)^\top$ has

$$\begin{aligned}\|(0.6, 0.8)^\top\|^2 &= 0.36 + 0.64 \\ &= 1,\end{aligned}$$

so it is in the closed unit ball but not the open one.

- For $p = 1$, the open ball of radius r around c is the open interval $(c - r, c + r)$.

Definition 0.12 (Second partial derivative and mixed partial derivative). Let f be a scalar-valued function of a $p \times 1$ vector $\tilde{x}$, and let $i, j \in \{1, \dots, p\}$. Suppose the partial derivative $\frac{\partial}{\partial x_j} f$ (Definition 0.1) exists at every point of an open ball around $\tilde{x}$ (Definition 0.11). The **second partial derivative** of f at $\tilde{x}$, first with respect to x_j and then with respect to x_i , is

$$\frac{\partial}{\partial x_i} \left(\frac{\partial}{\partial x_j} f(\tilde{x}) \right),$$

the partial derivative with respect to x_i of the function $\frac{\partial}{\partial x_j} f$, when it exists. When $i \neq j$, it is a **mixed partial derivative**.

Example 0.24 (Second partial derivatives of $x_1^2 x_2 + 3x_2$). For $f(\tilde{x}) = x_1^2 x_2 + 3x_2$ of Example 0.1, $\frac{\partial}{\partial x_1} f(\tilde{x}) = 2x_1 x_2$ and $\frac{\partial}{\partial x_2} f(\tilde{x}) = x_1^2 + 3$. Differentiating each of these again:

•

$$\begin{aligned}\frac{\partial}{\partial x_1} \left(\frac{\partial}{\partial x_1} f(\tilde{x}) \right) &= \frac{\partial}{\partial x_1} (2x_1 x_2) \\ &= 2x_2;\end{aligned}$$

•

$$\begin{aligned}\frac{\partial}{\partial x_2} \left(\frac{\partial}{\partial x_1} f(\tilde{x}) \right) &= \frac{\partial}{\partial x_2} (2x_1 x_2) \\ &= 2x_1;\end{aligned}$$

•

$$\begin{aligned}\frac{\partial}{\partial x_1} \left(\frac{\partial}{\partial x_2} f(\tilde{x}) \right) &= \frac{\partial}{\partial x_1} (x_1^2 + 3) \\ &= 2x_1;\end{aligned}$$

•

$$\begin{aligned}\frac{\partial}{\partial x_2} \left(\frac{\partial}{\partial x_2} f(\tilde{x}) \right) &= \frac{\partial}{\partial x_2} (x_1^2 + 3) \\ &= 0.\end{aligned}$$

The two mixed partial derivatives are both $2x_1$.

Definition 0.13 (Hessian matrix). Let f be a scalar-valued function of a $p \times 1$ vector $\tilde{x}$ whose first partial derivatives exist on an open ball around $\tilde{x}$ (Definition 0.11) and whose second partial derivatives (Definition 0.12) exist at $\tilde{x}$. The **Hessian matrix** of f at $\tilde{x}$ is the derivative (Definition 0.6) of the transposed gradient $(\frac{\partial}{\partial \tilde{x}} f(\tilde{x}))^\top$ (Definition 0.2):

$$\underbrace{\mathbf{H}_f(\tilde{x})}_{p \times p} \stackrel{\text{def}}{=} \frac{\partial}{\partial \tilde{x}} \left(\frac{\partial}{\partial \tilde{x}} f(\tilde{x}) \right)^\top. \quad (4)$$

By Theorem 0.1, the transposed gradient is the row-vector derivative $\frac{\partial f(\tilde{x})}{\partial \tilde{x}^\top}$ (Definition 0.5), so the Hessian can also be written

$$\begin{aligned}\mathbf{H}_f(\tilde{x}) &= \frac{\partial}{\partial \tilde{x}} \left(\frac{\partial f(\tilde{x})}{\partial \tilde{x}^\top} \right) \\ &= \frac{\partial^2 f(\tilde{x})}{\partial \tilde{x} \partial \tilde{x}^\top}.\end{aligned} \quad (5)$$

Both forms follow the shape rule of matrix multiplication (Remark 0.1): the $p \times 1$ column-vector operator $\frac{\partial}{\partial \tilde{x}}$ stands on the left of a $1 \times p$ row vector, giving a $p \times p$ matrix. By Definition 0.6, with $y_j = \frac{\partial}{\partial x_j} f(\tilde{x})$, the (i, j) entry of the Hessian is

$$[\mathbf{H}_f(\tilde{x})]_{ij} = \frac{\partial}{\partial x_i} \left(\frac{\partial}{\partial x_j} f(\tilde{x}) \right).$$

Example 0.25 (A Hessian matrix). Let $f(\tilde{x}) = e^{2x_1+x_2} - x_1$ for $\tilde{x} = (x_1, x_2)^\top$, where e is Euler's number, and write $u = 2x_1 + x_2$. By the chain rule, $\frac{\partial}{\partial x_1} e^u = 2e^u$ and $\frac{\partial}{\partial x_2} e^u = e^u$, so

$$\frac{\partial}{\partial \tilde{x}} f(\tilde{x}) = \begin{bmatrix} 2e^u - 1 \\ e^u \end{bmatrix}.$$

Differentiating each entry again,

$$\begin{aligned}\mathbf{H}_f(\tilde{x}) &= \begin{bmatrix} \frac{\partial}{\partial x_1}(2e^u - 1) & \frac{\partial}{\partial x_1}e^u \\ \frac{\partial}{\partial x_2}(2e^u - 1) & \frac{\partial}{\partial x_2}e^u \end{bmatrix} \\ &= \begin{bmatrix} 4e^u & 2e^u \\ 2e^u & e^u \end{bmatrix},\end{aligned}$$

and at $\tilde{x} = \tilde{0}$, where $u = 0$, $\mathbf{H}_f(\tilde{0}) = \begin{bmatrix} 4 & 2 \\ 2 & 1 \end{bmatrix}$.

Example 0.26 (A function with a gradient but no Hessian at a point). Let $f(\tilde{x}) = x_1|x_1|$ for $\tilde{x} = (x_1, x_2)^\top$, where $|\cdot|$ is the [absolute value](#). For $x_1 > 0$, $f = x_1^2$ and $\frac{\partial}{\partial x_1}f = 2x_1$; for $x_1 < 0$, $f = -x_1^2$ and $\frac{\partial}{\partial x_1}f = -2x_1$; and at $x_1 = 0$ the [difference quotient](#) is $h|h|/h = |h| \rightarrow 0$. So the gradient exists everywhere: it is $(2|x_1|, 0)^\top$. But $2|x_1|$ has no derivative in x_1 at $x_1 = 0$ (its difference quotient $2|h|/h$ is 2 for $h > 0$ and -2 for $h < 0$), so $[\mathbf{H}_f(\tilde{x})]_{11}$, and with it the Hessian, does not exist at any $\tilde{x}$ with $x_1 = 0$.

Videos

Jon Krohn’s “Calculus for Machine Learning” YouTube playlist has videos on second and higher partial derivatives:

- [Higher-Order Partial Derivatives](#)
- [Exercise on Higher-Order Partial Derivatives](#)

Definition 0.14 (Continuity in several variables). A function f from $\mathbb{R}^p$ to $\mathbb{R}^q$ is **continuous at** $\tilde{x}$ if for every $\varepsilon > 0$ there is a $\delta > 0$ such that $\|f(\tilde{y}) - f(\tilde{x})\| < \varepsilon$ whenever $\|\tilde{y} - \tilde{x}\| < \delta$. It is **continuous on** a set if it is continuous at every point of the set. For $q = 1$, $\|f(\tilde{y}) - f(\tilde{x})\| = |f(\tilde{y}) - f(\tilde{x})|$. For

$$\begin{aligned}p &= q \\ &= 1,\end{aligned}$$

this definition is the usual [continuity](#), written out with the ε - δ definition of the limit $\lim_{y \rightarrow x} f(y) = f(x)$.

Continuity survives the usual operations:

- A [composition](#) of continuous functions is continuous: choose the δ for the [outer function](#) first, and use it as the ε for the inner one.
- A sum $f + g$ of continuous real-valued functions is continuous, since $|(f + g)(\tilde{y}) - (f + g)(\tilde{x})| \leq |f(\tilde{y}) - f(\tilde{x})| + |g(\tilde{y}) - g(\tilde{x})|$: use $\varepsilon/2$ for each.

- A constant multiple cf is continuous, since $|cf(\tilde{y}) - cf(\tilde{x})| = |c| |f(\tilde{y}) - f(\tilde{x})|$: use $\varepsilon/(|c| + 1)$ for f .

Example 0.27 (A continuous function, and a discontinuous one).

- $f(\tilde{x}) = x_1 + x_2$ is continuous at every $\tilde{x}$. Each $|y_i - x_i| \leq \|\tilde{y} - \tilde{x}\|$, since the squared length is a sum of nonnegative squares, so

$$\begin{aligned} |f(\tilde{y}) - f(\tilde{x})| &= |(y_1 - x_1) + (y_2 - x_2)| && \text{(subtract)} \\ &\leq |y_1 - x_1| + |y_2 - x_2| && \text{(triangle inequality for numbers)} \\ &\leq 2 \|\tilde{y} - \tilde{x}\|, && \text{(each term is at most } \|\tilde{y} - \tilde{x}\|) \end{aligned}$$

which is less than ε whenever $\|\tilde{y} - \tilde{x}\| < \delta = \varepsilon/2$.

- $f(\tilde{x}) = 1$ if $x_1 > 0$ and $f(\tilde{x}) = 0$ otherwise is not continuous at $\tilde{0}$: for $\varepsilon = \frac{1}{2}$ and any $\delta > 0$, the point $\tilde{y} = (\delta/2, 0)^\top$ has $\|\tilde{y} - \tilde{0}\| = \delta/2 < \delta$ but $|f(\tilde{y}) - f(\tilde{0})| = 1 > \frac{1}{2}$.

Example 0.28 (Partial derivatives can exist where a function is not continuous). Let $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ be

$$f(\tilde{x}) = \begin{cases} \frac{x_1 x_2}{x_1^2 + x_2^2} & \text{if } \tilde{x} \neq \tilde{0}, \\ 0 & \text{if } \tilde{x} = \tilde{0}. \end{cases}$$

Both partial derivatives exist at $\tilde{0}$. For $h \neq 0$,

$$\begin{aligned} f(h, 0) &= \frac{h \cdot 0}{h^2 + 0} \\ &= 0, \end{aligned}$$

so

$$\begin{aligned} \frac{f(0 + h, 0) - f(0, 0)}{h} &= \frac{0 - 0}{h} && (f(h, 0) = 0 \text{ and } f(0, 0) = 0) \\ &= 0, && (0/h = 0 \text{ for } h \neq 0) \end{aligned}$$

and the limit as $h \rightarrow 0$ is 0: $\frac{\partial}{\partial x_1} f(\tilde{0}) = 0$ (Definition 0.1). The same steps with $f(0, h) = 0$ give $\frac{\partial}{\partial x_2} f(\tilde{0}) = 0$.

But f is not continuous at $\tilde{0}$ (Definition 0.14). Take $\varepsilon = \frac{1}{4}$ and any $\delta > 0$, and let $t = \delta/2$ and $\tilde{y} = (t, t)^\top$. Then

$$\begin{aligned}
\|\tilde{y} - \tilde{0}\| &= \sqrt{t^2 + t^2} \\
&= t\sqrt{2} \\
&= \delta/\sqrt{2} \\
&< \delta,
\end{aligned}$$

but

$$\begin{aligned}
|f(\tilde{y}) - f(\tilde{0})| &= \left| \frac{t \cdot t}{t^2 + t^2} - 0 \right| \quad (\tilde{y} \neq \tilde{0}, \text{ since } t > 0) \\
&= \frac{t^2}{2t^2} \quad (t^2 + t^2 = 2t^2) \\
&= \frac{1}{2}, \quad (\text{cancel } t^2 > 0)
\end{aligned}$$

which is more than $\varepsilon = \frac{1}{4}$.

The partial derivatives only look along the two coordinate axes, where f is 0; along the line $x_1 = x_2$, f is $\frac{1}{2}$ everywhere except at $\tilde{0}$. So partial derivatives existing at a point says little about how f behaves near it.

Theorem 0.10 (Symmetry of the Hessian). *If the second partial derivatives of f exist and are continuous (Definition 0.14) on an open ball around $\tilde{x}$ (Definition 0.11), then $\frac{\partial}{\partial x_i} \left(\frac{\partial}{\partial x_j} f(\tilde{x}) \right) = \frac{\partial}{\partial x_j} \left(\frac{\partial}{\partial x_i} f(\tilde{x}) \right)$ for all i, j , so $\mathbf{H}_f(\tilde{x})$ (Definition 0.13) is symmetric (symmetric matrix).*

The proof applies the one-variable mean value theorem twice, which these notes do not develop; see (Rudin 1976, Theorem 9.41), which is stated for two variables: apply it to f as a function of x_i and x_j , with the other coordinates held fixed.

Example 0.29 (Mixed partial derivatives agree). In Example 0.25, the $(1, 2)$ and $(2, 1)$ entries of $\mathbf{H}_f(\tilde{x})$ are both $2e^{2x_1+x_2}$.

Example 0.30 (Without continuity, the mixed partials can differ). Let $f(x_1, x_2) = \frac{x_1 x_2 (x_1^2 - x_2^2)}{x_1^2 + x_2^2}$ for $\tilde{x} \neq \tilde{0}$, and $f(\tilde{0}) = 0$. For $x_2 \neq 0$, the partial derivative (Definition 0.1) in x_1 at $(0, x_2)$ is

$$\begin{aligned}
\frac{\partial}{\partial x_1} f(0, x_2) &= \lim_{h \rightarrow 0} \frac{f(h, x_2) - f(0, x_2)}{h} && \text{(definition of the partial derivative, Definition 0.1)} \\
&= \lim_{h \rightarrow 0} \frac{x_2(h^2 - x_2^2)}{h^2 + x_2^2} && (f(0, x_2) = 0; \text{cancel } h) \\
&= \frac{x_2 \cdot (-x_2^2)}{x_2^2} \\
&= -x_2, && \text{(the quotient is continuous at } h = 0)
\end{aligned}$$

and

$$\begin{aligned}
\frac{\partial}{\partial x_1} f(\tilde{0}) &= \lim_{h \rightarrow 0} (0 - 0)/h \\
&= 0,
\end{aligned}$$

so $\frac{\partial}{\partial x_1} f(0, x_2) = -x_2$ holds at $x_2 = 0$ too. In the same way, with the roles of x_1 and x_2 swapped, $f(x_1, k)/k = x_1(x_1^2 - k^2)/(x_1^2 + k^2) \rightarrow x_1$, so $\frac{\partial}{\partial x_2} f(x_1, 0) = x_1$ for every x_1 . So at $\tilde{0}$ (Definition 0.13)

$$\begin{aligned}
[\mathbf{H}_f(\tilde{0})]_{21} &= \frac{\partial}{\partial x_2} \left(\frac{\partial}{\partial x_1} f \right) \\
&= \frac{d}{dx_2} (-x_2) \\
&= -1, \\
[\mathbf{H}_f(\tilde{0})]_{12} &= \frac{\partial}{\partial x_1} \left(\frac{\partial}{\partial x_2} f \right) \\
&= \frac{d}{dx_1} x_1 \\
&= 1 :
\end{aligned}$$

the Hessian at $\tilde{0}$ is not symmetric. By Theorem 0.10, then, the second partial derivatives of this f cannot all be continuous near $\tilde{0}$.

Theorem 0.11 (Hessian of a quadratic form). *If $\mathbf{S}$ is a symmetric $p \times p$ matrix that is constant with respect to $\tilde{x}$, then $f(\tilde{x}) = \tilde{x}^\top \mathbf{S} \tilde{x}$ has $\mathbf{H}_f(\tilde{x}) = 2\mathbf{S}$ for every $\tilde{x}$.*

Proof

Proof.

$$\begin{aligned}
 \mathbf{H}_f(\tilde{x}) &= \frac{\partial}{\partial \tilde{x}} \left(\frac{\partial}{\partial \tilde{x}} (\tilde{x}^\top \mathbf{S} \tilde{x}) \right)^\top && \text{(Definition 0.13)} \\
 &= \frac{\partial}{\partial \tilde{x}} (2\mathbf{S}\tilde{x})^\top && \text{(Theorem 0.7)} \\
 &= \left(\frac{\partial}{\partial \tilde{x}} \tilde{x} \right) (2\mathbf{S})^\top && \left(\frac{\partial}{\partial \tilde{x}} (2\mathbf{S}\tilde{x})^\top = \frac{\partial}{\partial \tilde{x}} (2\mathbf{S}\tilde{x}), \text{ Definition 0.6} \right) \\
 &= \mathbf{I}_p (2\mathbf{S})^\top && \left(\frac{\partial}{\partial \tilde{x}} \tilde{x} = \mathbf{I}_p, \text{ Remark 0.3} \right) \\
 &= 2\mathbf{S}. && (\mathbf{S} \text{ is symmetric})
 \end{aligned}$$

□

Example 0.31 (The Hessian of a 2×2 quadratic form). For $\mathbf{S} = \begin{bmatrix} 3 & 1 \\ 1 & 2 \end{bmatrix}$ as in Example 0.16, Theorem 0.11 gives $\mathbf{H}_f(\tilde{x}) = \begin{bmatrix} 6 & 2 \\ 2 & 4 \end{bmatrix}$. Directly, the gradient found there, with β_i renamed x_i , is $(6x_1 + 2x_2, 2x_1 + 4x_2)^\top$, and differentiating its entries by x_1 and by x_2 (Definition 0.13) gives

$$\begin{aligned}
 \mathbf{H}_f(\tilde{x}) &= \begin{bmatrix} \frac{\partial}{\partial x_1}(6x_1 + 2x_2) & \frac{\partial}{\partial x_1}(2x_1 + 4x_2) \\ \frac{\partial}{\partial x_2}(6x_1 + 2x_2) & \frac{\partial}{\partial x_2}(2x_1 + 4x_2) \end{bmatrix} \\
 &= \begin{bmatrix} 6 & 2 \\ 2 & 4 \end{bmatrix}.
 \end{aligned}$$

Further reading

- Marsden and Tromba (2013) is a standard textbook on multivariable and vector calculus. It covers differentiation of functions of several variables, multiple integrals, line and surface integrals, and the theorems of Green, Gauss, and Stokes.
- Petersen and Pedersen (2012) is a free desk reference that collects matrix identities. Its chapter on derivatives lists derivatives of vector and matrix expressions, such as the linear and quadratic forms on this page.
- Miller (2016) The partial derivative review exercises and worked solutions on this page are adapted from this supplemental review chapter.

See also the [Linear Algebra and Vector Calculus further reading](#).

- Hua Zhou’s lecture notes for “UCLA Biostat 216 - Mathematical Methods for Biostatistics” (2023 Fall)
- [Neural Networks](#) is a YouTube playlist by Grant Sanderson (3Blue1Brown); its chapters on gradient descent and backpropagation calculus visually illustrate how gradients of multivariate cost functions are computed and used for optimization.

References

- Felippa, Carlos A. n.d. *Matrix Calculus*. Appendix F of Introduction to Finite Element Methods, course notes, University of Colorado Boulder. Accessed October 6, 2026. https://quickfem.com/wp-content/uploads/IFEM.AppF_.pdf.
- Fieller, Nick. 2016. *Basics of Matrix Algebra for Statistics with R*. Chapman; Hall/CRC. <https://doi.org/10.1201/9781315370200>.
- Hutchinson, Brian. n.d. *DATA 471/571 (Machine Learning) and CSCI 481/581 (Deep Learning) Video Lectures*. Western Washington University. Accessed September 28, 2026. https://facultyweb.cs.wvu.edu/~hutchib2/video_lectures/data371/.
- Marsden, Jerrold E., and Anthony Tromba. 2013. *Vector Calculus*. 6th ed. Macmillan Learning. <https://www.macmillanlearning.com/college/us/product/Vector-Calculus/p/1429215089>.
- Miller, Steven J. 2016. *The Probability Lifesaver: Calculus Review Problems*. https://web.williams.edu/Mathematics/sjmiller/public_html/probabilitylifesaver/index.htm#:~:text=http%3A//web.williams.edu/Mathematics/sjmiller/public_html/probabilitylifesaver/supplementalchap_calcreview.pdf.
- Petersen, Kaare Brandt, and Michael Syskind Pedersen. 2012. *The Matrix Cookbook*. Technical University of Denmark. <https://www2.imm.dtu.dk/pubdb/edoc/imm3274.pdf>.
- Rudin, Walter. 1976. *Principles of Mathematical Analysis*. 3rd ed. International Series in Pure and Applied Mathematics. McGraw-Hill.
- Wikipedia contributors. n.d. *Gradient* — *Wikipedia, the Free Encyclopedia*. <https://en.wikipedia.org/wiki/Gradient>.